

From a Statement of Work to a plan you can sign

An internal platform that turns a signed client contract into a costed, scheduled, staffed project plan — and proves, line by line, where every part of it came from.

PREPARED FOR	Board of Directors
SUBJECT	SOW-to-Project Execution Platform — status, evidence and recommendation
STAGE	Working pilot, verified end to end on a live client contract
REPOSITORY	github.com/Muhannad-Khaled/Tasks-Management-System
EVIDENCE	NileBank consumer loyalty programme — 7-page contract, processed end to end

How to read this report. Section 4 contains the measured results from a real client contract. Section 6 states the platform’s known limits without softening them — a system built to declare what it does not know should be described the same way.

Two technical diagrams accompany this document: the component architecture and the processing pipeline.

01 Executive summary

Every client engagement begins with the same manual step: someone reads twenty pages of contract prose and rebuilds it by hand into tasks, dates, owners and acceptance criteria. It costs days of senior time, it is done differently by every project manager, and the link back to the contract is lost the moment the plan is written.

This platform performs that step and keeps the link. It reads the signed contract, produces a complete project plan, and attaches to every statement in that plan a reference to the passage of the contract that produced it. Where the contract is silent, it says so and raises the question with the client rather than quietly filling the gap.

It has been run end to end on a live client contract. The results are in section 4.

Drafting the plan is the easy part. Knowing which parts the contract actually authorised is the product.

What the board is being asked to note

- The platform works and has produced a verified plan from a real contract.
- Its distinguishing feature is not speed. It is that its output can be audited against the contract months later, by anyone.
- It is a pilot: one team, no access control, and AI capacity on a trial allowance. Section 7 sets out what production readiness requires.

02 The problem, and why general AI does not solve it

Any general-purpose AI tool will produce a polished project plan from any contract in thirty seconds. It will also invent a go-live date, a pricing rate and a named approver — in the same tone of voice as the parts it got right. There is no way to tell them apart.

That is not a cosmetic problem. A plan whose origins cannot be checked is a plan that cannot be signed off, invoiced against, or defended in a scope dispute. The commercial risk is not that the system is slow; it is that it is confidently wrong in a way nobody notices until delivery.

Three costs the manual process carries today

Senior time	Days on every engagement, before a single hour of delivery work begins.
Lost traceability	Six weeks into delivery, nobody can answer which clause put a given task on the board.
Hidden assumptions	The contract never named a launch date, so someone picked one — and it entered the plan looking exactly like a fact.

03 What the platform produces

One contract goes in. What comes out is a working plan, published to the delivery board the team already uses.

OUTPUT	WHAT IT CONTAINS
Task breakdown	The work split across commercial, technical and operations, each task tied to the requirement it delivers.
Schedule	Sequenced by dependency, with the critical path identified and the float on every other task made explicit.
Owners	Every task assigned to the role that does it, and staffed from the people directory where a candidate exists.
Acceptance criteria and tests	What “done” means for each piece of work, using the measurable thresholds the contract actually stated.
Register of assumptions	Every working value supplied because the contract did not state one — recorded and visible, never presented as fact.
Client questions	The same gaps, turned into a list to send back before they become change requests.

The distinction the platform is built around

Every value in the plan is marked with how it came to exist. This single distinction is what separates the output from a document that merely sounds convincing.

STATED The contract says it outright. Carries a reference to the exact passage.	INFERRED It follows from something the contract does state. Carries the passage it follows from.	ASSUMED The contract is silent; a working value was supplied so planning could continue. Becomes a recorded assumption and a client question.
--	---	--

Verification is per statement, not per plan

A task is not graded once and waved through. It is broken into the individual assertions it makes — *this takes fifteen hours, this is the account manager’s responsibility, this cannot start before signature* — and each is checked against the contract on its own. A task can be entirely right about the work and quietly wrong about the duration; a single overall score hides exactly the detail a project manager needs to catch.

04 Evidence: one contract, end to end

The figures below are the recorded output of a single run against the NileBank consumer loyalty contract — a seven-page document of 2,188 words. They are read directly from the platform’s database, not estimated.

11	812	100%	9
tasks scheduled	hours estimated	statements verified	client questions raised
PRODUCED	COUNT	NOTE	
Requirements identified	15	each referenced to the contract section that states it	
Tasks, scheduled and sequenced	11	across commercial (2), technical (5) and operations (4)	
Roles defined	11	work assigned by role; the directory fills names where it can	
User stories	20	15 written for the client, 5 written for the engineer doing the work	
Acceptance criteria	43	8 carry a measured threshold taken from the contract	
Test cases	22	2 flagged because they rest on an assumed value	
Assumptions → client questions	9 → 9	raised before delivery rather than discovered during it	
Statements checked against the contract	35	every one supported by cited text	

What was checked, and how

The 35 statements were extracted from the eleven tasks and each judged against the passage it cites. A sample, verbatim from the record:

STATEMENT	CONTRACT SAYS
Accrual rate is 2 points per 10 EGP	"2 points for every 10 EGP of qualifying spend"
Hypercare support is provided for 30 days	"hypercare support for 30 calendar days"
Pilot runs across 5 branches for one week	"5 nominated branches for a period of one week"
Authentication uses OAuth 2.0 with mTLS	stated in the technical requirements

No statement cited a passage that does not exist, and none was marked supported without a citation. Both are checked mechanically, because a fabricated reference would otherwise be indistinguishable from a real one.

ON A PERFECT SCORE

A previous run of the same contract scored 90%, and its weakest task scored 33% — it had conflated two separate deliverables and asserted a figure the contract applies to something else. That task went to the board carrying a `GROUNDING-FAILED` label, and the project manager approved it knowing that. The check is not decorative: it has caught real errors, and the platform does not hide them.

05 Control, and what the platform refuses to do

The platform never publishes on its own. It prepares the plan, puts every assumption, weak task and unanswered question in front of the project manager, and waits for a named approval.

Before approval, the manager can edit any task, regenerate a weak one, answer a client question, or reject the plan outright; every action is recorded. After approval, cards are updated rather than duplicated, and if someone changes or deletes a card downstream the manager is told the plan and the board have diverged.

The refusals

- ✗ **It refuses a document it cannot genuinely read.** A scanned contract stops the process rather than yielding an empty plan that looks complete.
- ✗ **It refuses to report success on an empty result.** If nothing was extracted, it retries once and then reports failure.
- ✗ **It refuses to state a figure it cannot source.** An unsupported number in an acceptance criterion is dropped, not saved.
- ✗ **It refuses to invent people.** Work is assigned to roles; names come only from a real directory.
- ✗ **It refuses to fill a silence.** Where the contract specifies nothing, the card says so in plain words.

The refusals are the product. Anything can generate a plan; the commercial value is in a system that knows the boundary of what it was told.

06 Assurance and known limits

502

automated tests

9

independent checks per plan

39

planning gaps audited

100%

AI activity logged

Every request made to the AI, and the answer it gave, is stored against the project it served, so any output produced months ago can be traced to the instruction behind it. The scheduling, the validation and the audit trail are ordinary software: the AI is never asked to mark its own work. The full test suite runs in under three minutes without network access.

Limits, stated plainly

Roles, not named people	Work is assigned to a role, not an individual. There is no HR integration, and inventing staff names is the exact failure this platform exists to prevent. The project manager still makes the final allocation.
Coverage of the document is not yet verified	The platform proves that everything it produced traces back to the contract. It does not yet prove the reverse — that everything in the contract was picked up. On the test contract, 62% of passages are referenced by something in the plan. The gap is measurable, and closing it is the first item in section 7.
Output varies between runs	The same contract produced 5 requirements on one run and 15 on another. Both were verified against the document, but the shape of the plan is not yet stable. The cause is known: the platform never defines what one requirement is, so the model decides each time.
The delivery board is the weak link	The full plan lives in our own system; the board carries a summary. If structured tracking inside the board becomes a requirement, that is the piece to reconsider.
Built for a pilot, not a fleet	One team, no access control, and AI capacity on a trial allowance of twenty requests per day. All three are known work items before wider use.

07 Recommendation

The platform is ready to be used on live engagements under project-manager supervision, which is how it is designed to be used. Moving it beyond a single team requires the following, in order.

IMMEDIATE Prove document coverage	Report every passage of the contract that produced nothing, so an unexamined section and an examined one cannot look identical. The measurement already exists; it needs to become a check the manager sees before approving.
IMMEDIATE Process long contracts in the background	A large document currently holds up the upload, which has caused duplicated work. Moving it to a background job with live progress removes both.
NEAR TERM Stabilise the shape of the plan	Anchor a requirement to the contract's own clause structure so the count becomes a property of the document rather than of the run.

NEAR TERM Decline out-of-scope contracts loudly	A contract outside our service lines should be refused outright, not quietly turned into a plausible-looking plan.
---	--

INVESTMENT Commercial AI capacity	Removes the trial ceiling, cuts processing time substantially, and avoids being locked to a single supplier. Much of the engineering described above exists to work within twenty requests per day.
---	---

The value is not that it writes the plan. It is that it knows the difference between what the contract said and what it assumed — and tells you which is which.

Every competitor can produce a project plan from a document. What a client is buying here is the ability to open any task on the board, six months into delivery, and read the clause that authorised it.

SOW-to-Project Execution Platform · Board Report · September 2026

Source code and full history: github.com/Muhannad-Khaled/Tasks-Management-System

Figures in section 4 are read from the platform's own records for the NileBank engagement and can be reproduced from the repository.